

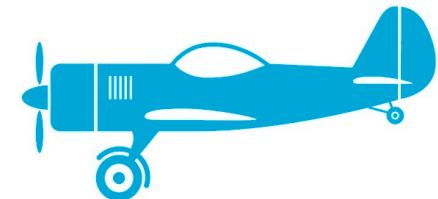
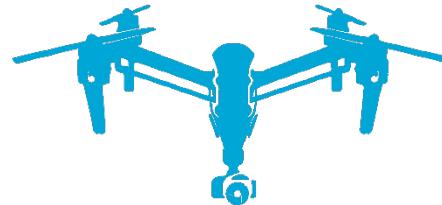
Deriving Timing Properties from System Traces Using Data-driven Techniques

Şerban Vădineanu

Supervisor: Mitra Nasri

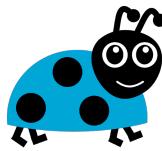
What systems?

- **Real-Time Systems (RTS):**
 - the system's correctness depends on both its **functional** and **temporal** correctness.



Why deriving properties?

Embedded systems' complexity is growing

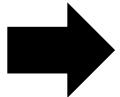


Complexity results in more bugs



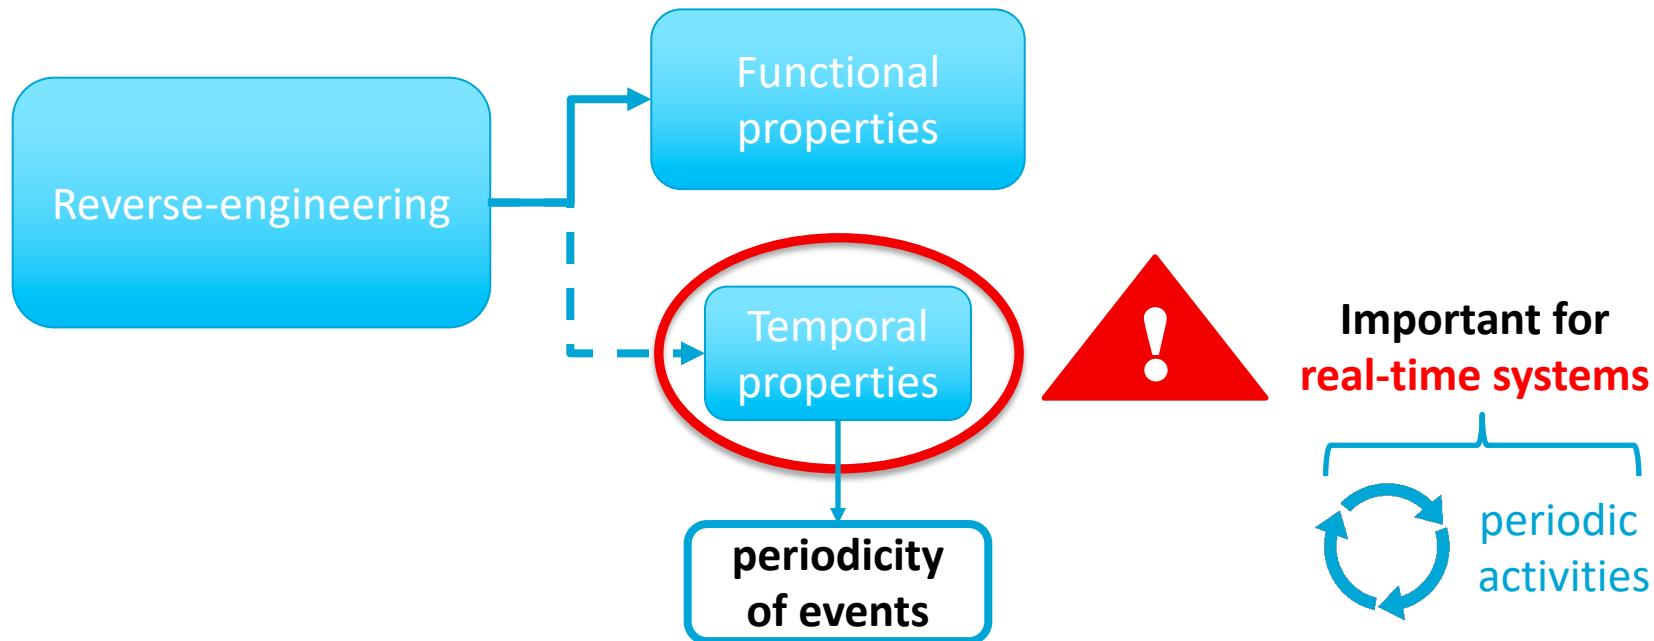
Complexity poses challenges on security

There is a need for tools to understand the system's **behavior**



Reverse-engineering

Which properties?



Why is period inference challenging?



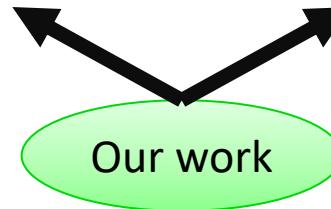
No operating system
Only one task



Lower-priority task
in a real-time operating system



High-priority aperiodic tasks,
sporadic tasks,
missed jobs,
release jitter



task—a functionality of the system, which periodically releases instances (**jobs**)



jitter—maximum deviation of the release time among all jobs of a task

What has been done?

Period estimation

Machine learning for
reverse-engineering
real-time systems

Reverse-engineering
timing properties of
real-time systems

- periodogram
 - autocorrelation
- } Fourier transform

- neural networks
- k-nearest neighbors

- only two works on period inference:

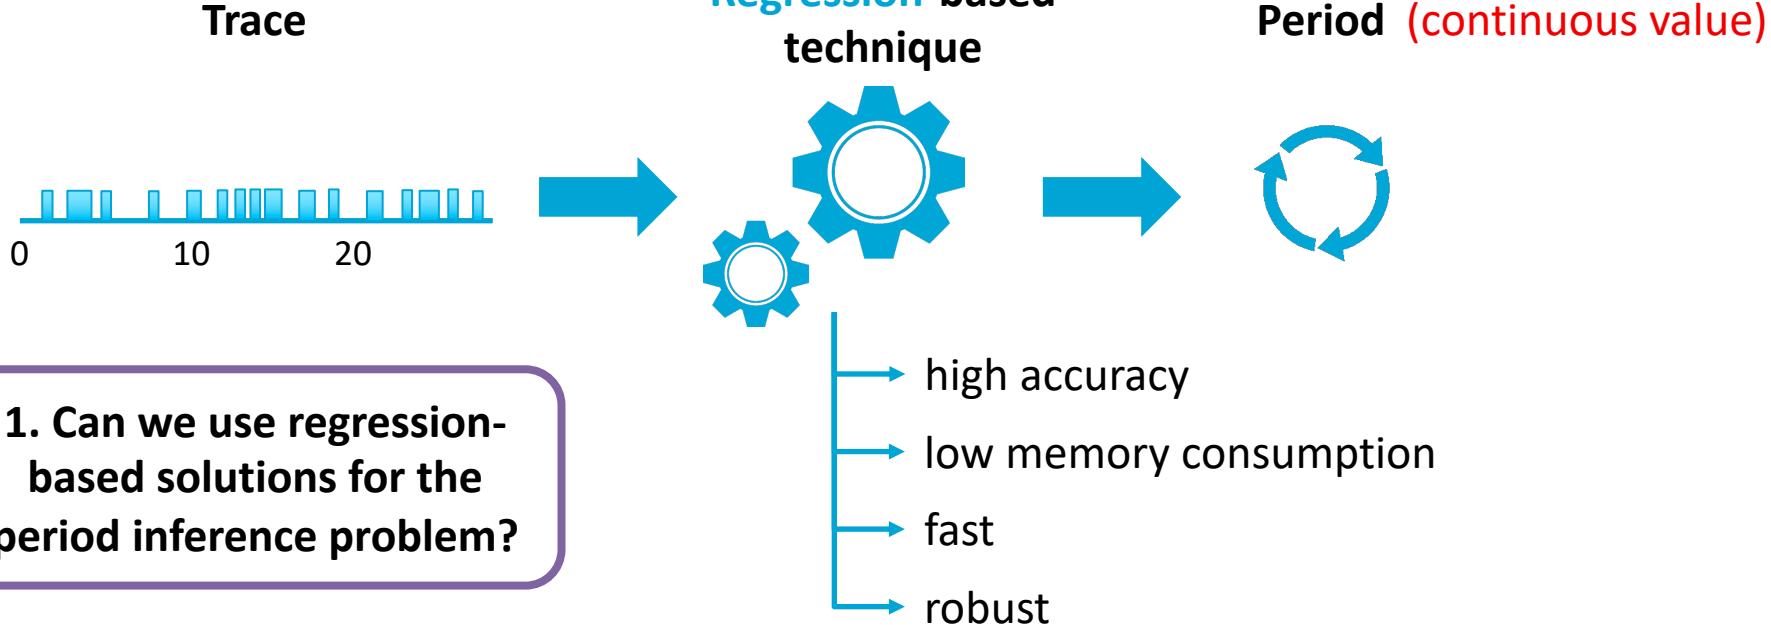
- PeTaMi
- Fourier transform-based

NO machine learning solution
for **period inference**

[1] O. Ilegorov, R. Torres, and S. Fischmeister. "Periodic task mining in embedded system traces". *Real-Time and Embedded Technology and Applications Symposium*, pp. 331-340, 2017.

[2] C. Young, H. Olufowobi, G. Bloom, and J. Zambreno. "Automotive intrusion detection based on constant CAN message frequencies across vehicle driving modes." *ACM Workshop on Automotive Cybersecurity*, pp. 9-14, 2019.

What is our aim?



robustness – the solution works even when the assumptions do not hold

- Period mining from real-time systems traces.
- Mines intervals between events and checks if they become periodic.
- Not learning based.

Pros ✓

1. **good performance** for harmonic periods
2. does **not require** additional learning

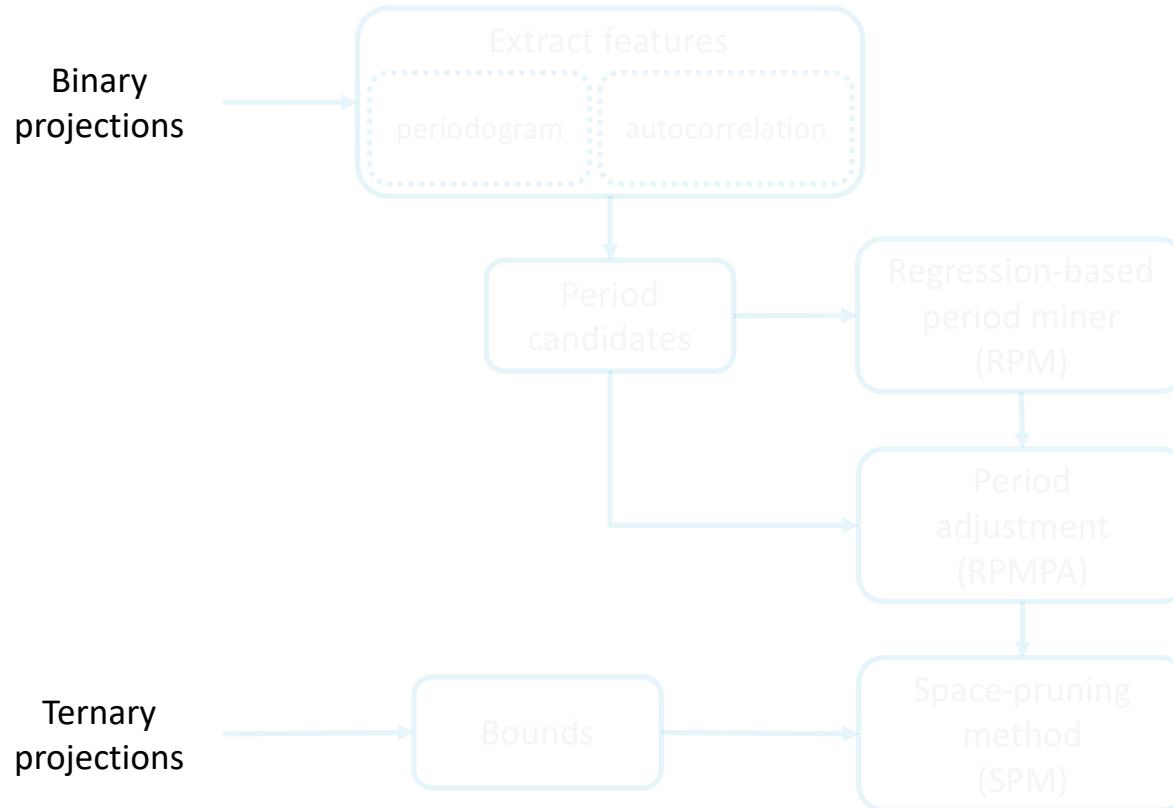
Cons ✗

1. **poor performance** for non-harmonic periods
2. **not robust** if there is execution time variation or release jitter
3. **slow** for longer traces



harmonic periods— every period divides all other smaller periods

Work in a nutshell



What is a projection?

Schedule

Binary projection
of task₂

Ternary projection
of task₂

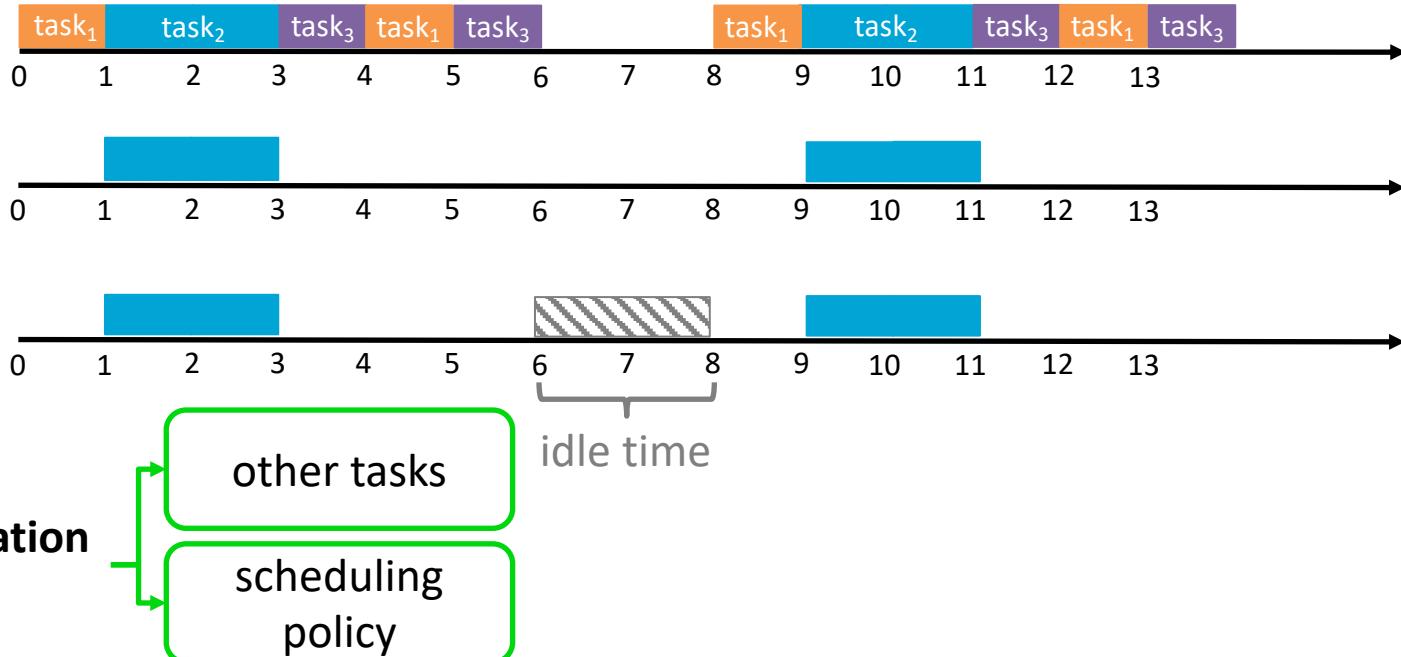
Input

No information
about

other tasks

scheduling
policy

idle time

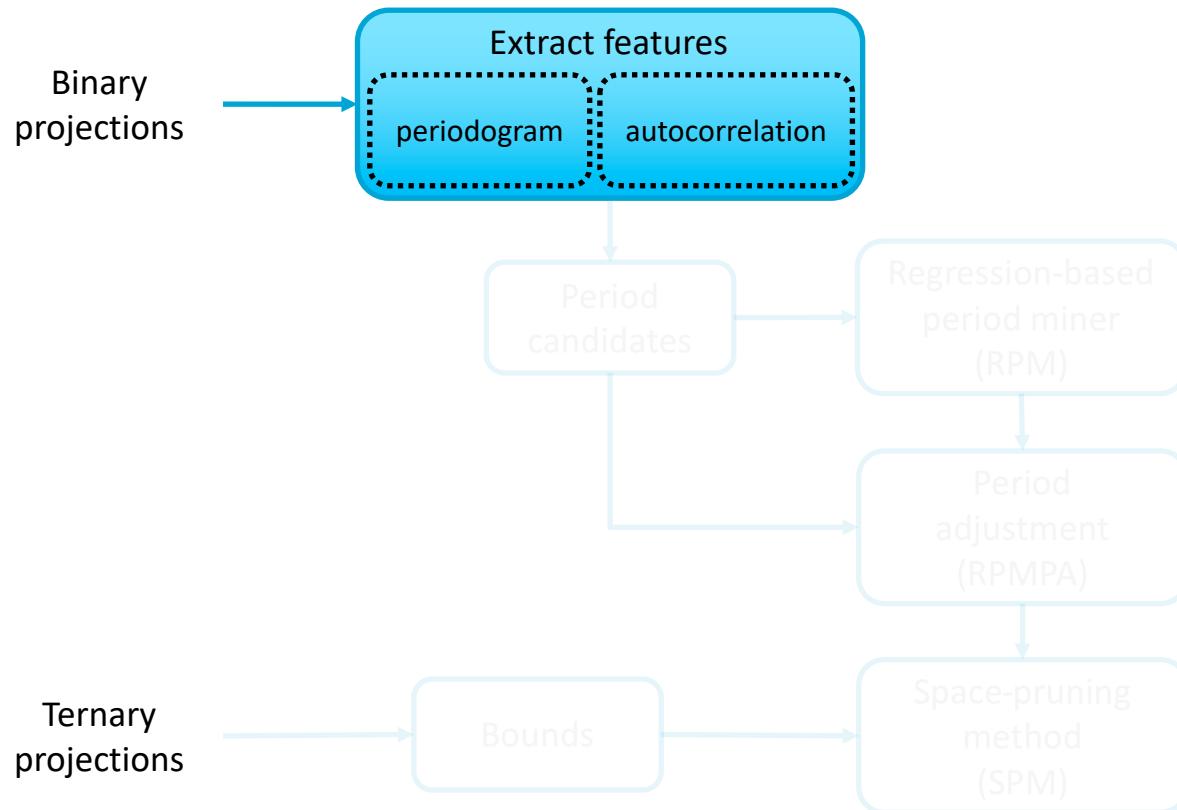


schedule— a particular assignment of tasks to the processor(s) and time intervals. It determines the task execution sequence.

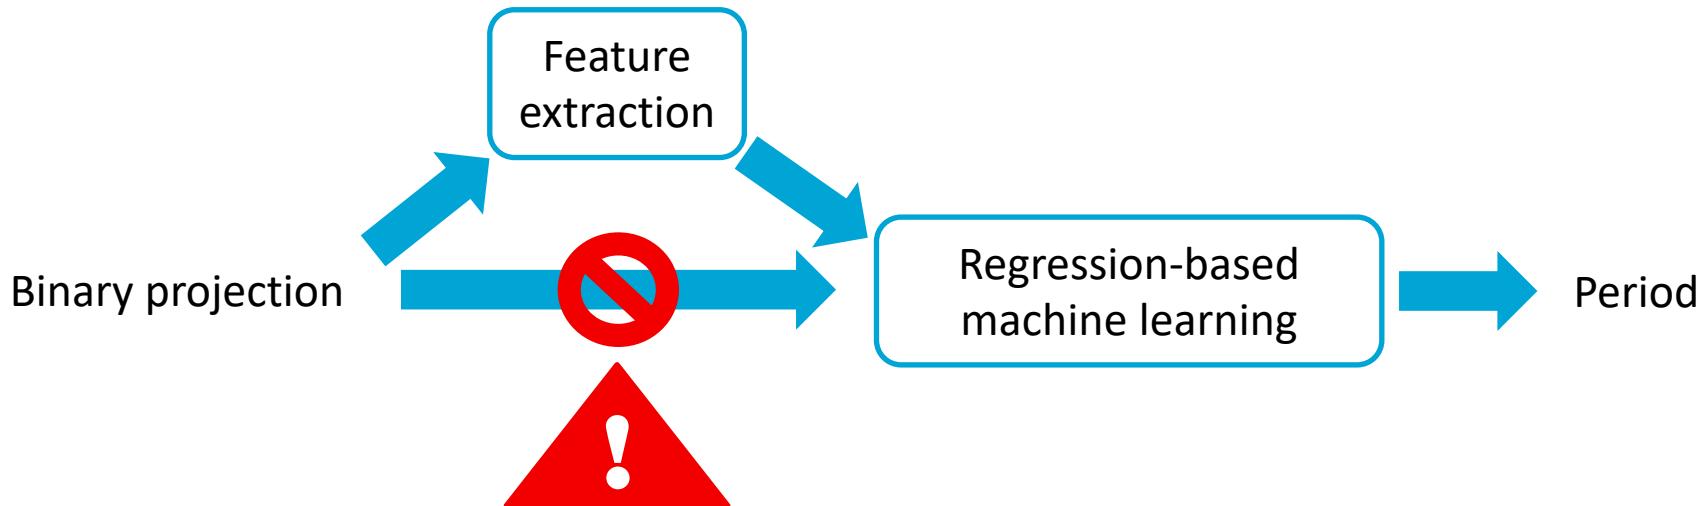


scheduling policy— the policy by which jobs are scheduled

Work in a nutshell

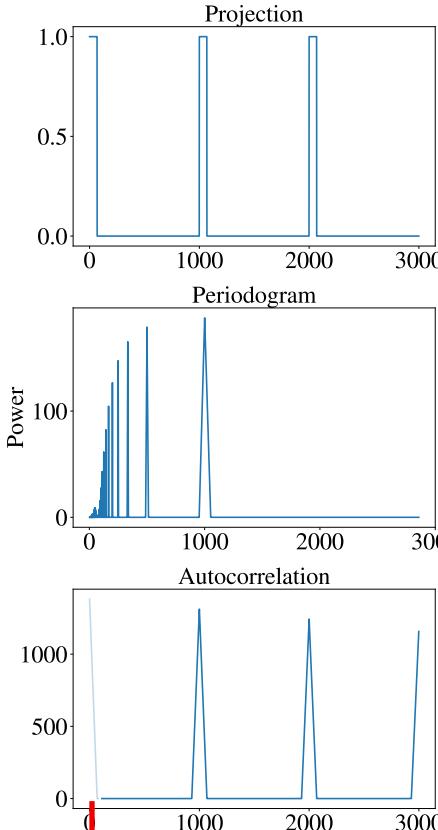


Why extracting features?



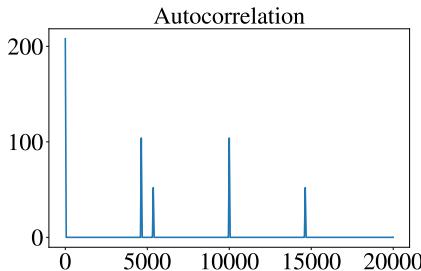
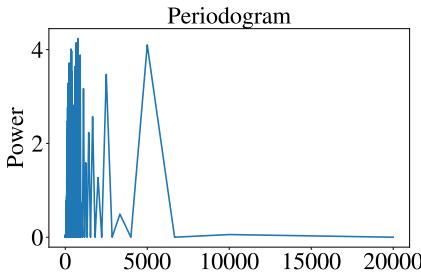
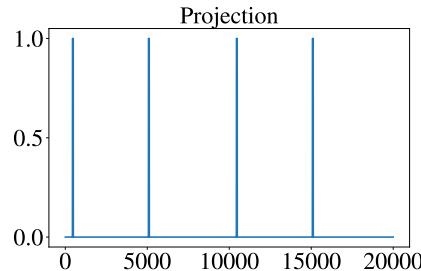
- very high dimensionality
- variable-length input

Feature extraction



- an estimate of the spectral density of the signal
- provided by the squared length of each Fourier coefficient of the signal
- examines how similar a sequence is to its previous values for different time lags
- inverse Fourier transform of dot product between the signal and its conjugation

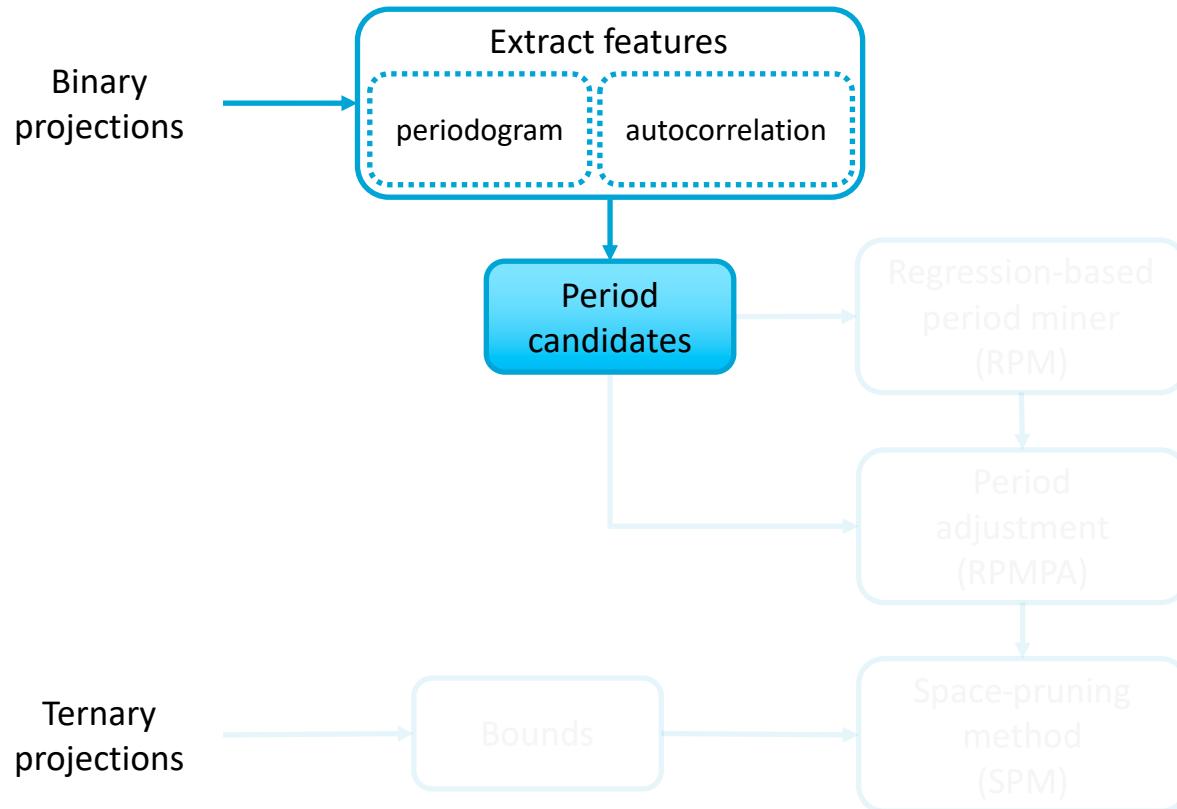
Challenges of feature extraction



- Have multiple peaks
- The highest peak is not necessarily the answer

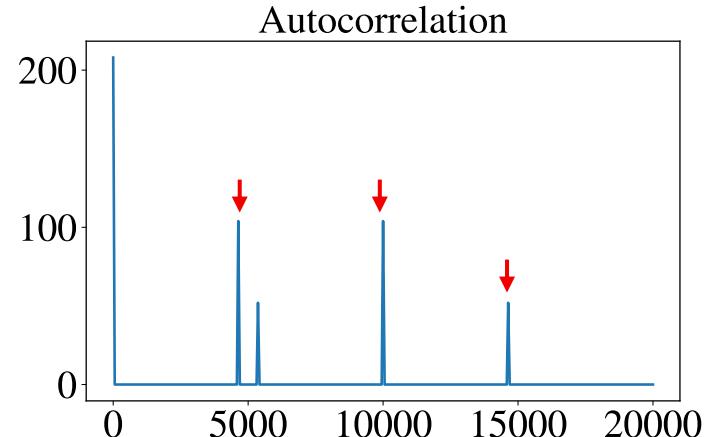
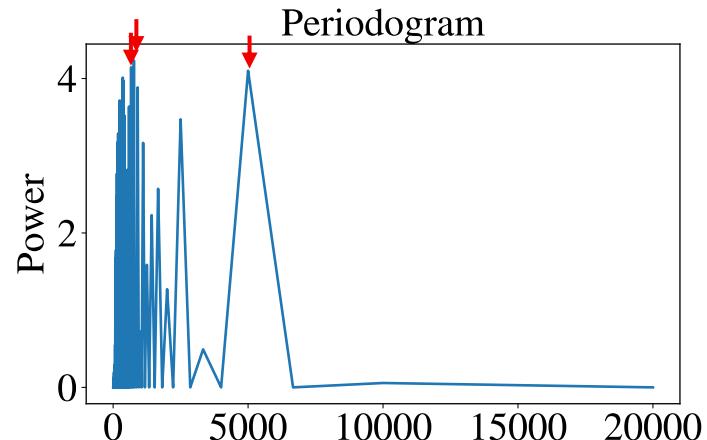
We decided to extract **period candidates** and employ regression models

Work in a nutshell

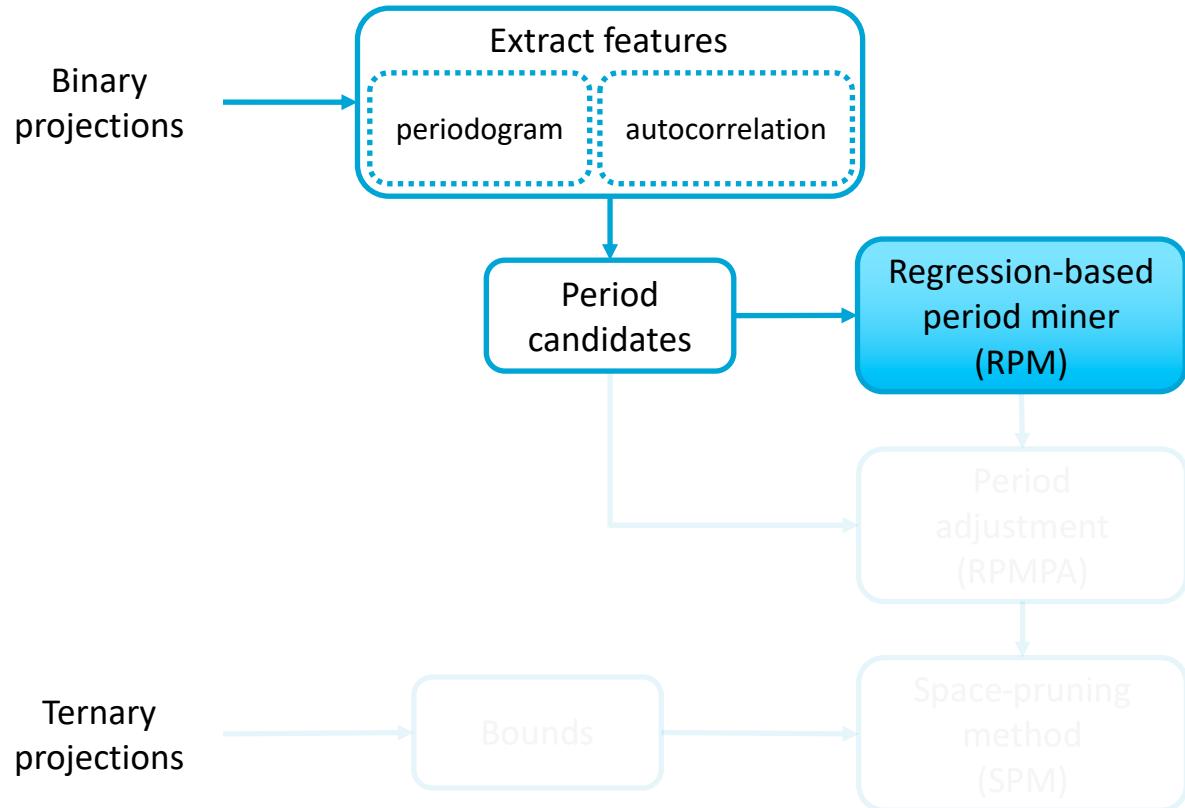


Period candidates

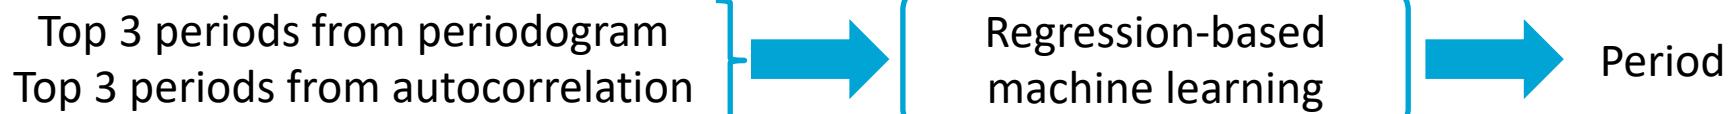
- top 20 periods from the **periodogram**
- top 20 periods from the **autocorrelation** as **candidates**
- top 3 periods from both as **features** for regression



Work in a nutshell



Regression-based period miner (RPM)



2. Which regression algorithm would result in a better accuracy for inferring periods?

Which regression algorithms?

cubist

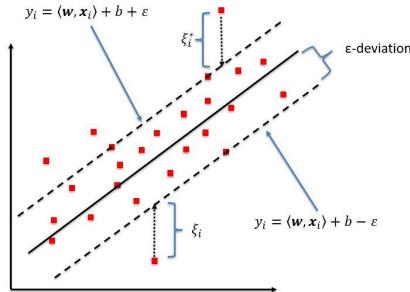
extremely randomized
regression trees
(extraTrees)

gradient boosting
(gbm)

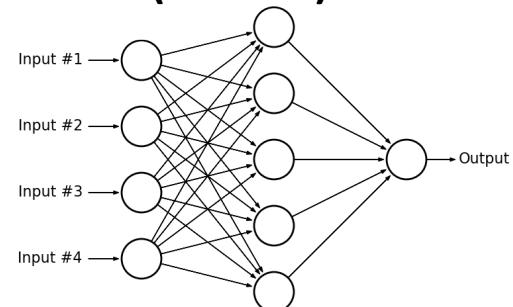
Bayesian additive
regression trees
(bartMachine)

Tree-based
algorithms

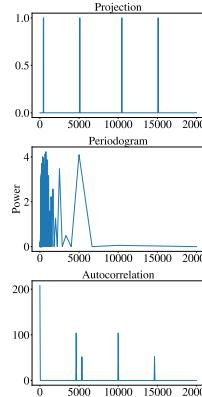
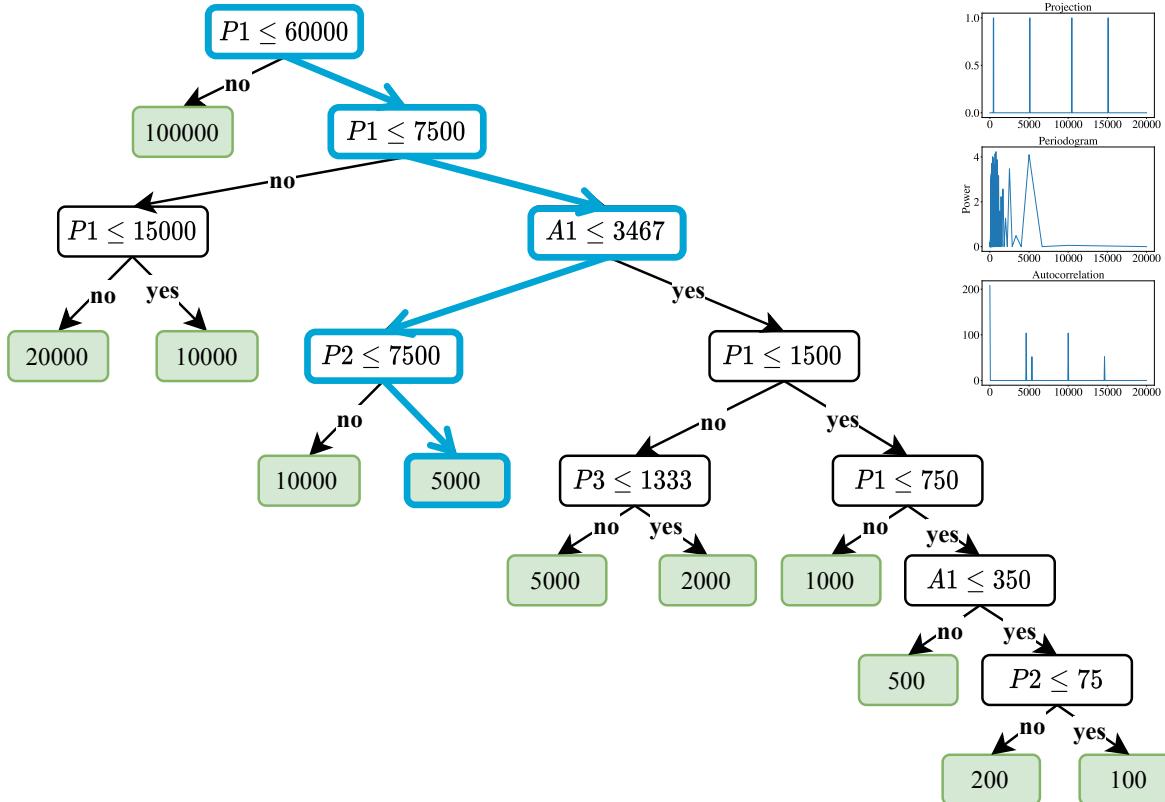
support vector regression
(svr)



averaged neural networks
(avNNet)



What is a regression tree?



period=5000

$P1=769, P2=666, P3=5000,$
 $A1=4635, A2=10000, A3=5365$

P1, P2, P3 – top 3 periods from periodogram
A1, A2, A3 – top 3 periods from autocorrelation

Which regression algorithms?

cubist

extremely randomized
regression trees

gradient boosting

Bayesian additive
regression trees



collapses the tree
into a set of rules



trains multiple trees
averages the output

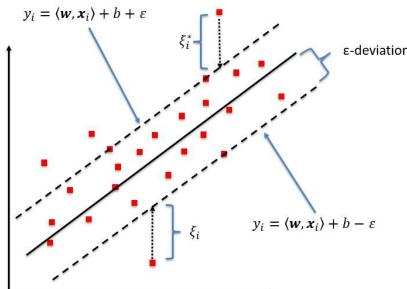


creates a tree to minimize
the error of previous trees

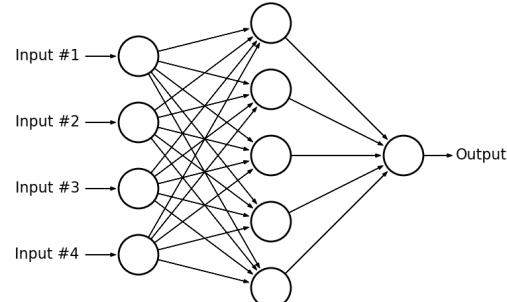


similar to gbm has prior
probabilities for tree size +
leaf values

support vector machine

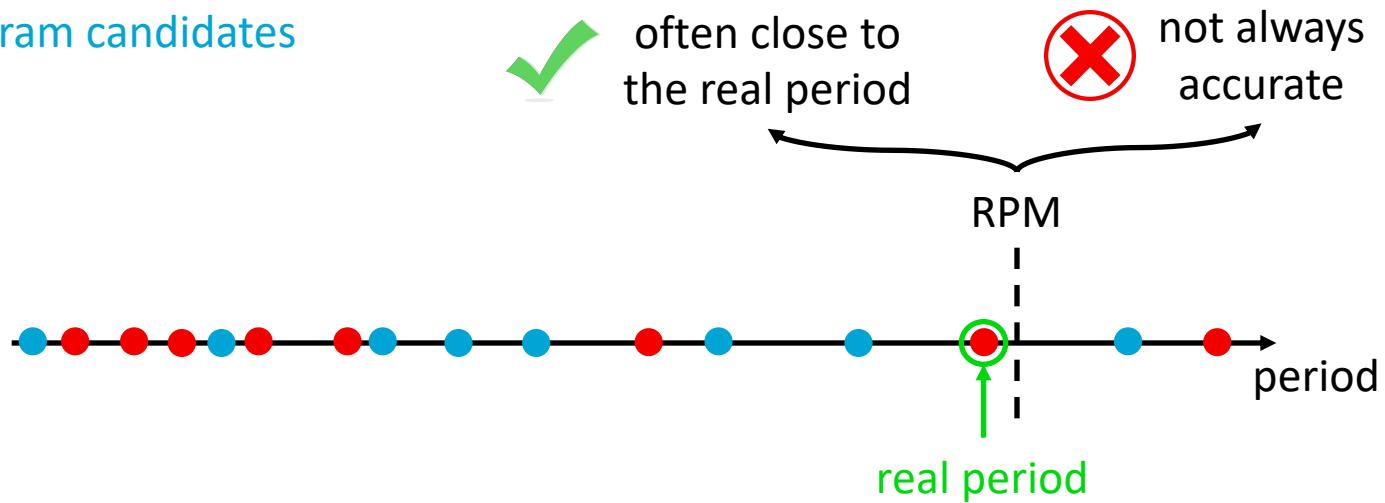


averaged neural networks

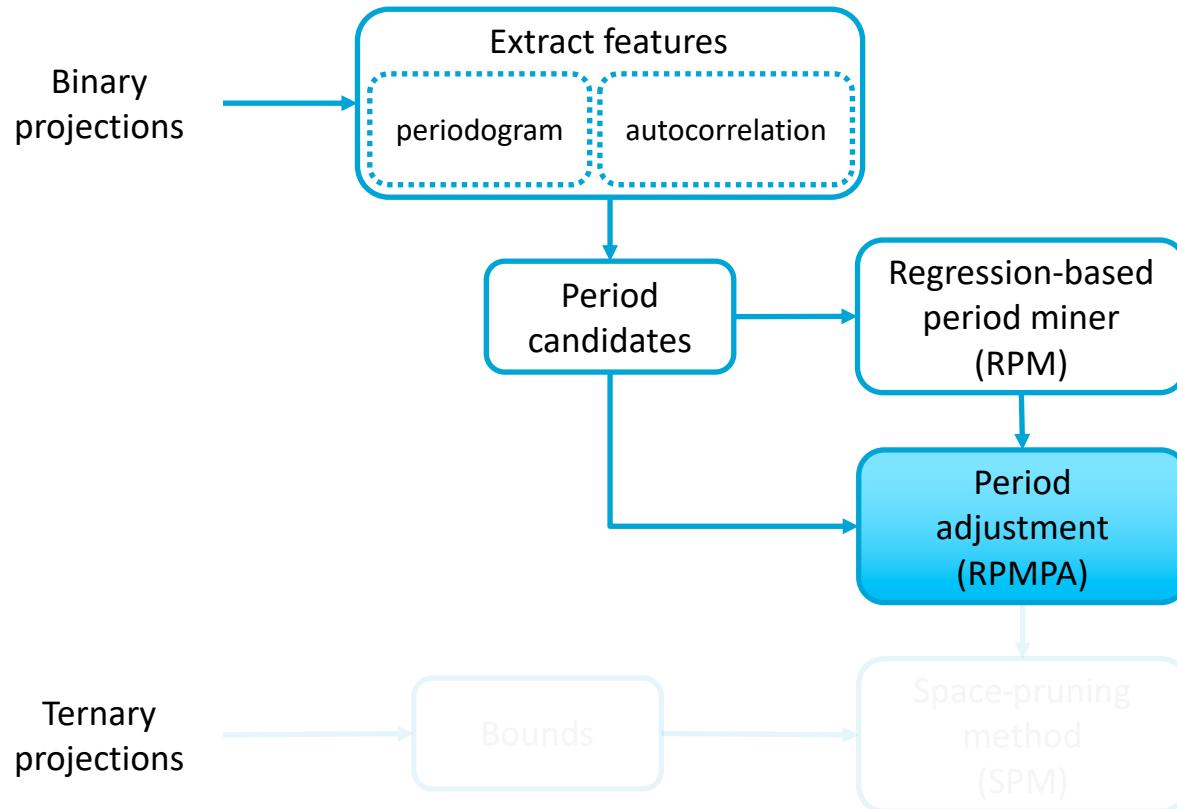


Observations

- autocorrelation candidates
- periodogram candidates

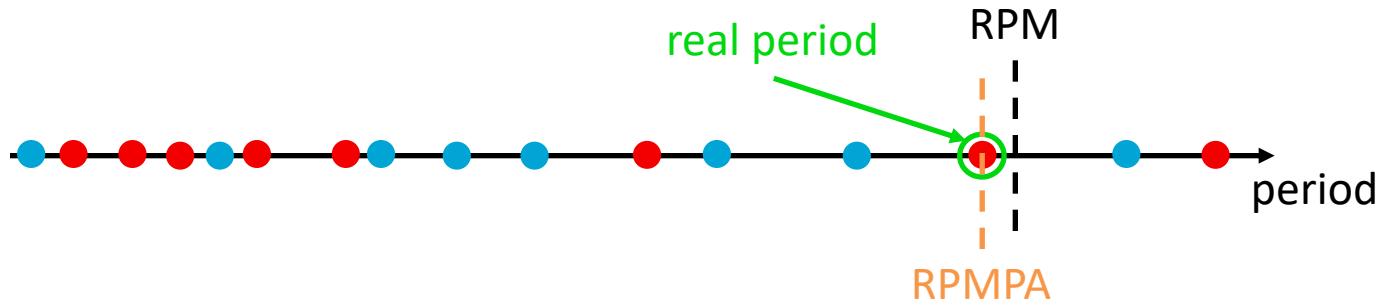


Work in a nutshell



Period adjustment (RPMPA)

1. Create candidate list from top 20 values of **periodogram** & **autocorrelation**
2. Choose the candidate closest to what RPM estimates

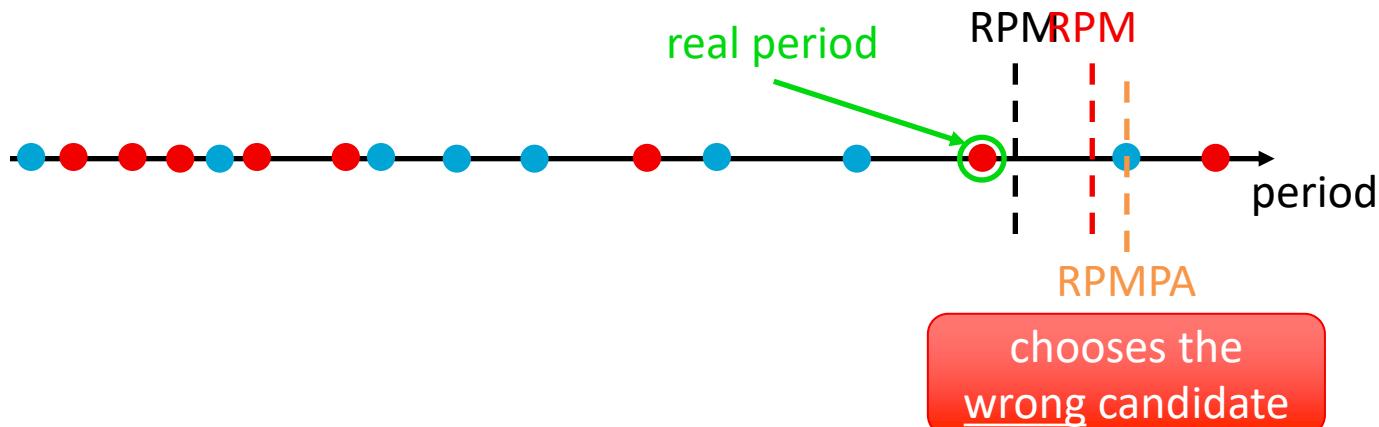


RPM – regression-based period miner

RPMPA – regression-based period miner with period adjustment

What can go wrong?

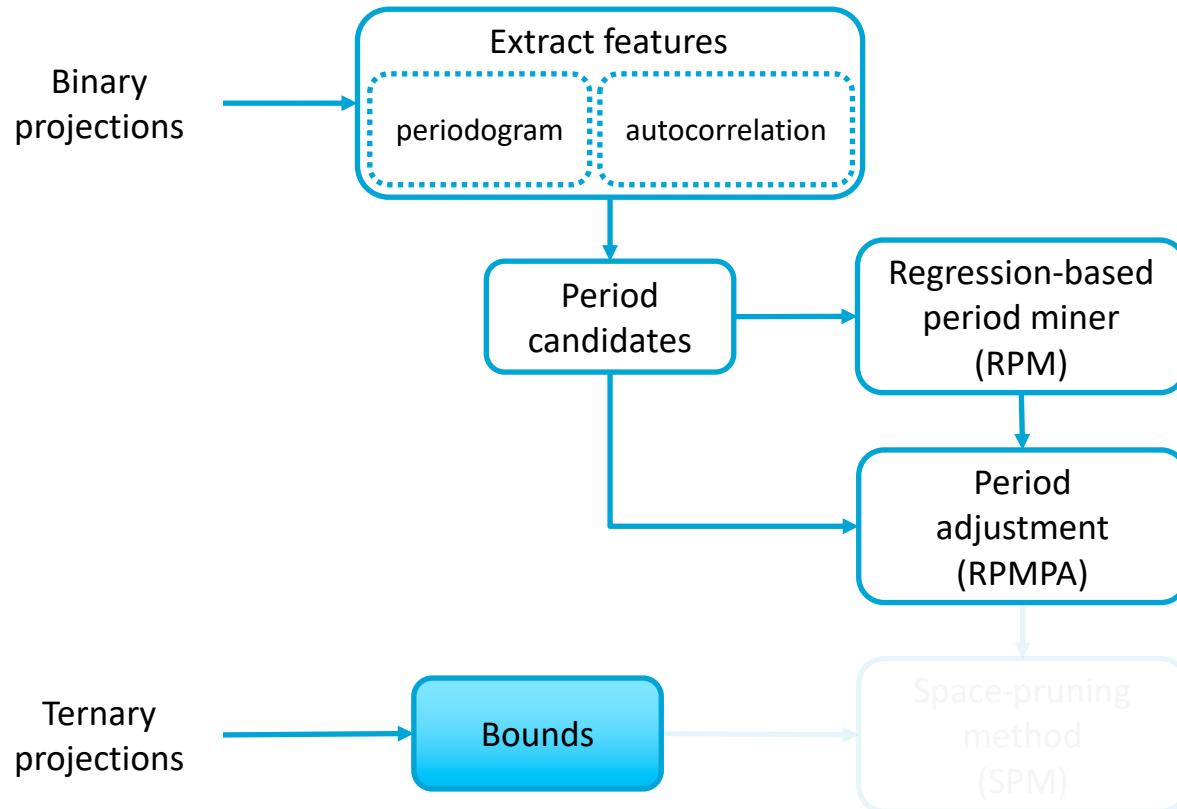
3. Can we derive a pair of bounds to restrict candidate values?



RPM – regression-based period miner

RPMPA – regression-based period miner with period adjustment

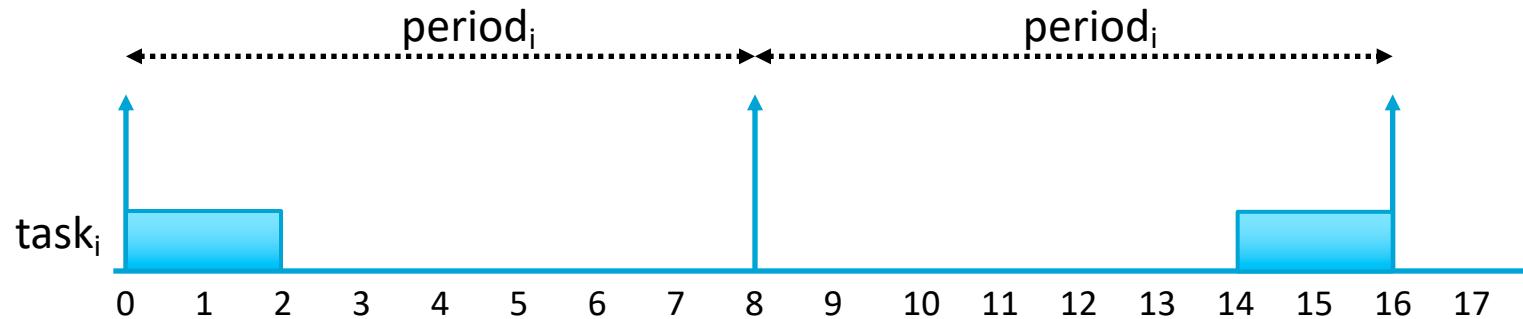
Work in a nutshell



Bounds on the period

1. **Lower bound** – based on longest non-running interval from a projection.
2. **Upper bound** – based on knowledge about the *idle time*.

Lower bound derivation

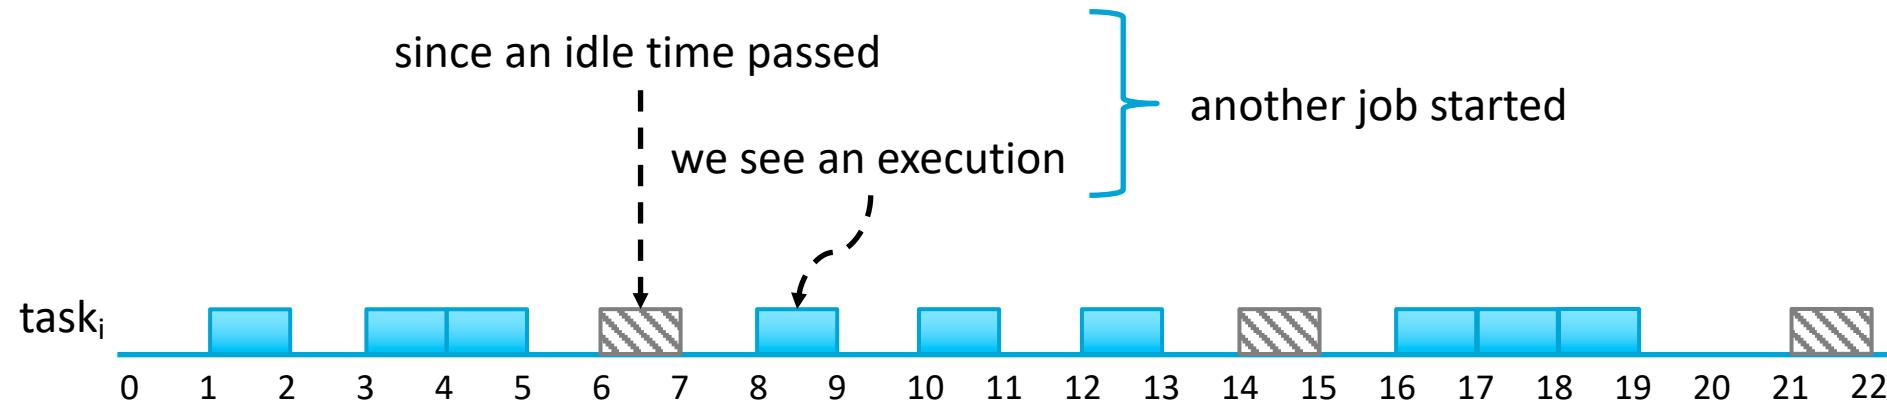


longest non-running interval
(LNI) $< 2 \times \text{period}_i$

period $_i = 8$
execution $_i = 2$
deadline $_i = \text{period}_i$

$$\text{period}_i > \frac{\text{LNI}}{2}$$

Upper bound derivation

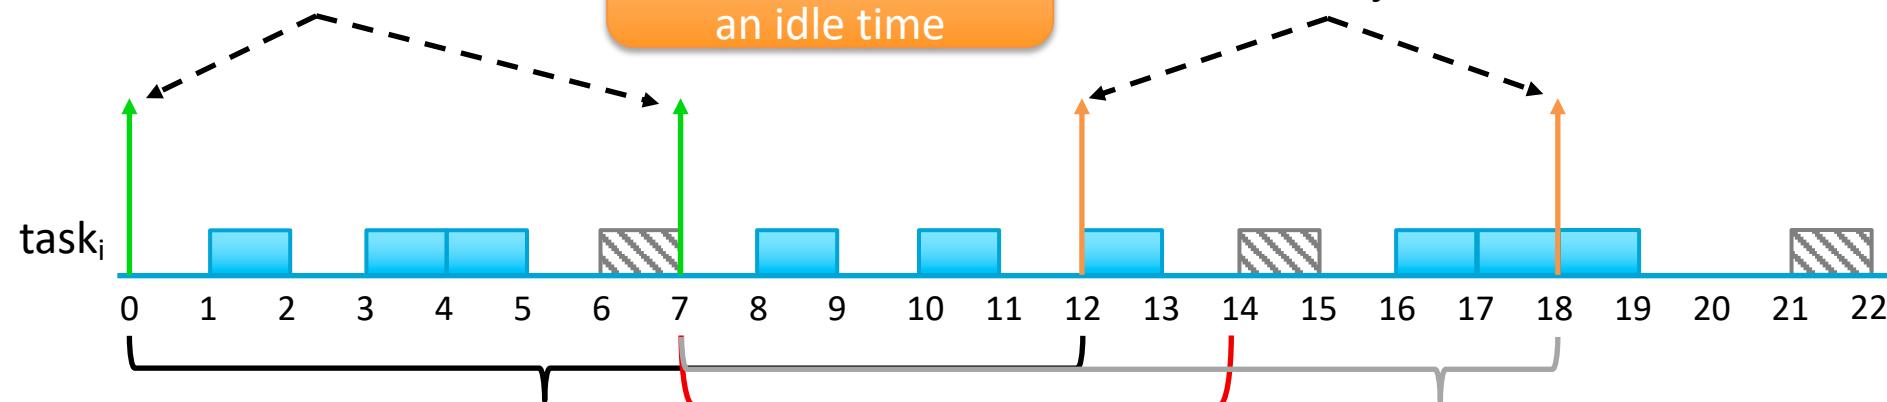


Upper bound derivation

earliest possible release of previous job

The last time unit of execution before an idle time

latest possible release of current job



The end of idle time

upper interval (UI) = 12

we don't know if here is just one job

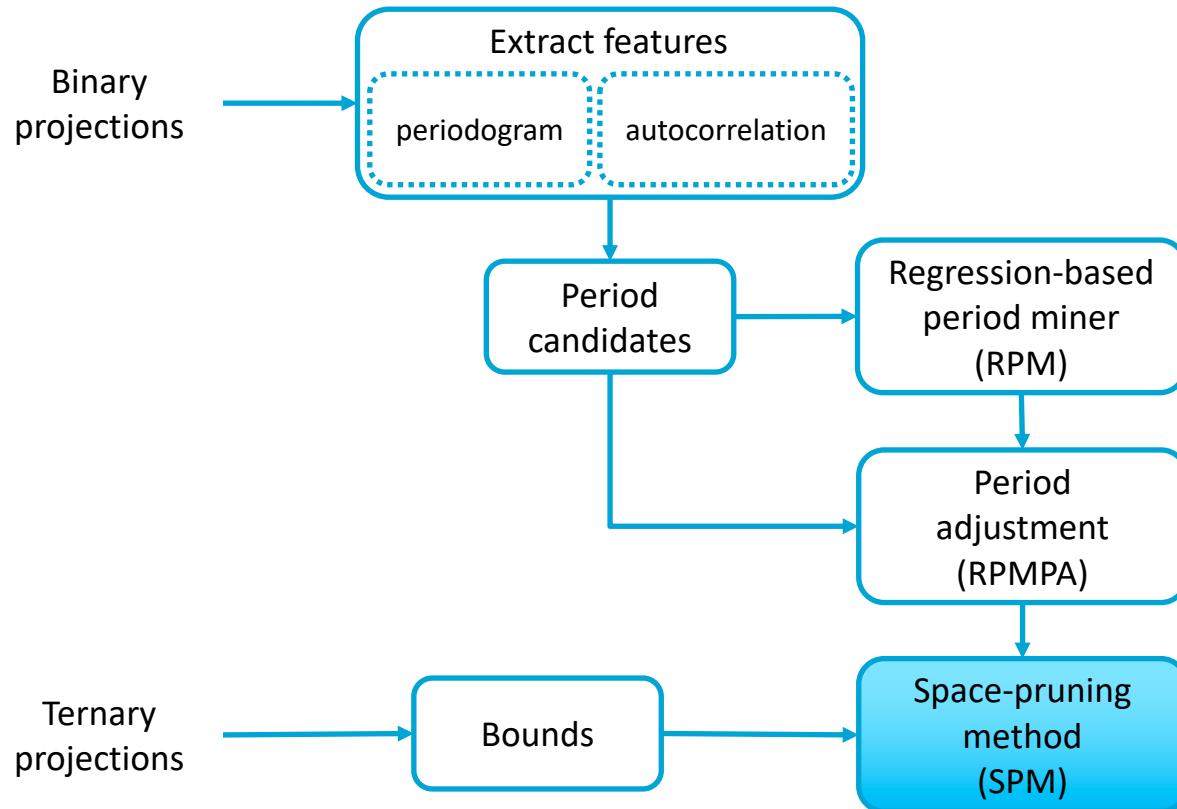
period = difference between the release time of consecutive jobs

$$\text{period}_i \leq \min(\text{UI})$$

Q: When is the release time?

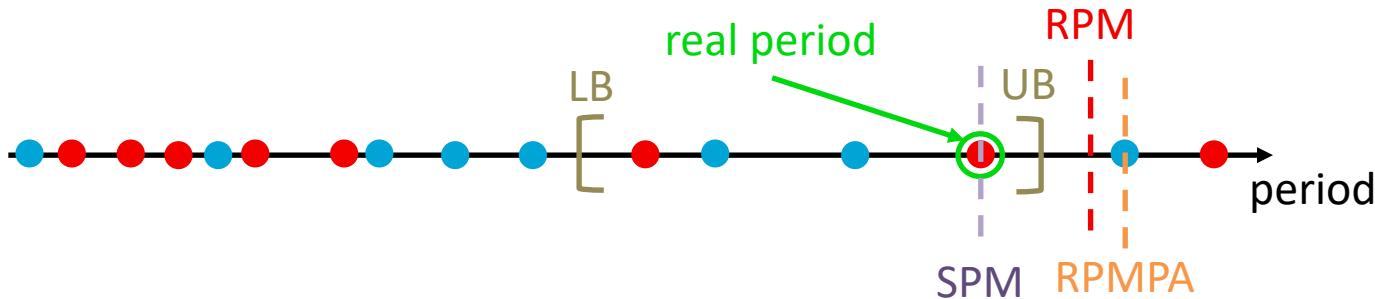
A: We consider extreme cases.

Work in a nutshell



Space pruning method (SPM)

1. Create candidate list from top 20 values of **periodogram** & **autocorrelation**
2. **Filter the candidates based on the bounds**
3. Choose the candidate closest to what RPM estimates



RPM – regression-based period miner

RPMPA – regression-based period miner with period adjustment

UB – upper bound

LB – lower bound

Experiments

What we want to find?

1. Do we have an improvement over the state of the art in terms of

Accuracy

Runtime

2. How do the six considered regression algorithms compare on the two aforementioned criteria?
3. How much does **SPM** improve the accuracy of **RPM** and **RPMPA**?
4. How do our solutions behave under non-ideal conditions?

RPM – regression-based period miner

RPMPA – regression-based period miner with period adjustment

SPM – space pruning method

Error metric

Average estimation error

$$e = \frac{\sum_{i=1}^N \frac{\widehat{T}_i - T_i}{T_i}}{N}$$

Estimated period

Actual periods

Number of tasks

Data set

Automotive benchmark application

- Task sets follow Autosar standard
- Periods from {1, 2, 5, 10, 20, 50, 100, 200, 1000}ms
- For simplicity: *automotive traces*

Synthetic task sets

- Random periods between 10ms and 1000ms with a base period of 10ms and log-uniform distribution (commonly used in papers)
- For simplicity: *log-uniform traces*

Simso simulator to generate logs and **extract** traces.

Experimental setups

We considered the effects of:

period ranges

utilization

scheduling
policy

release
jitter

high-priority
aperiodic tasks

missing
jobs

execution time
variation

tardiness

RPM

Data set: 2000 traces

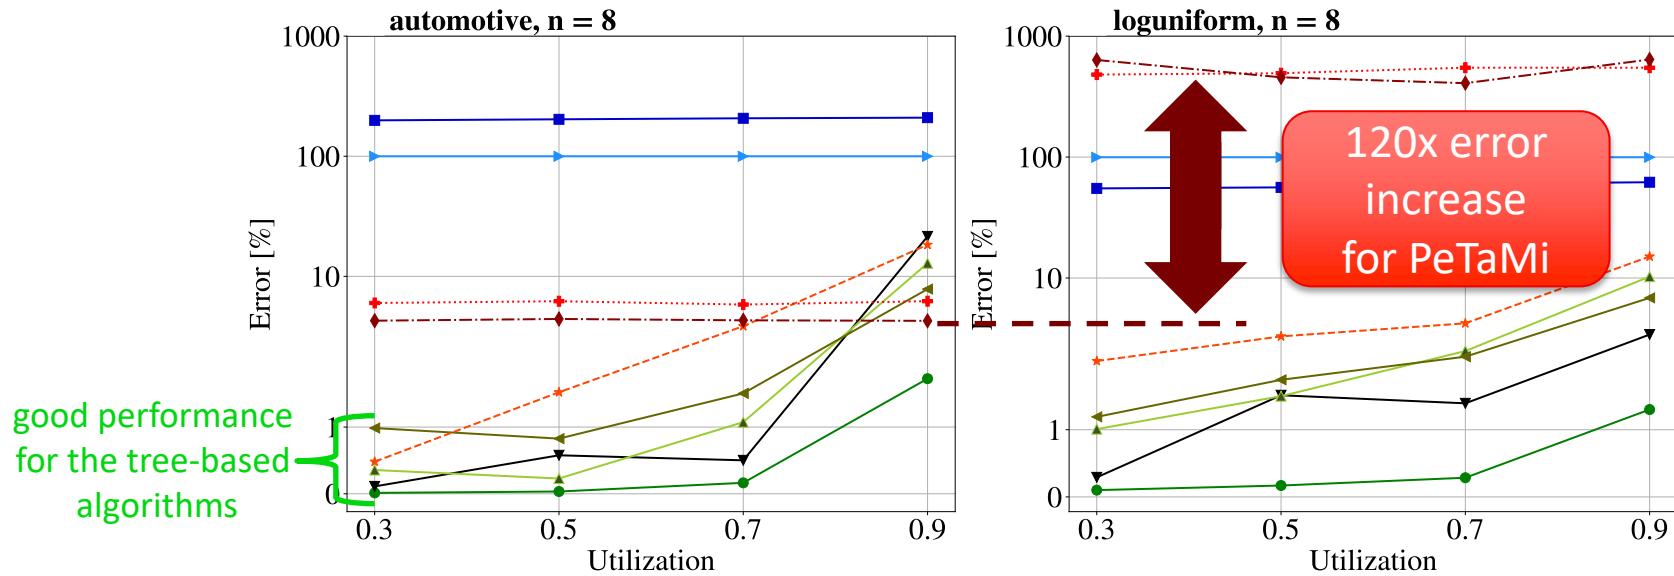
Error estimation: cross-validation

**4. How robust are our
solutions to sources of
non-determinism?**

Performance is dependent on the utilization

Legend:

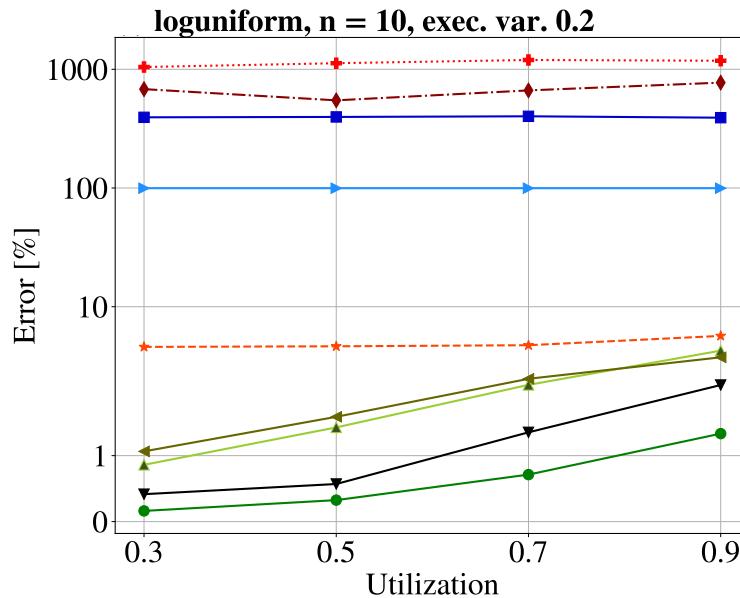
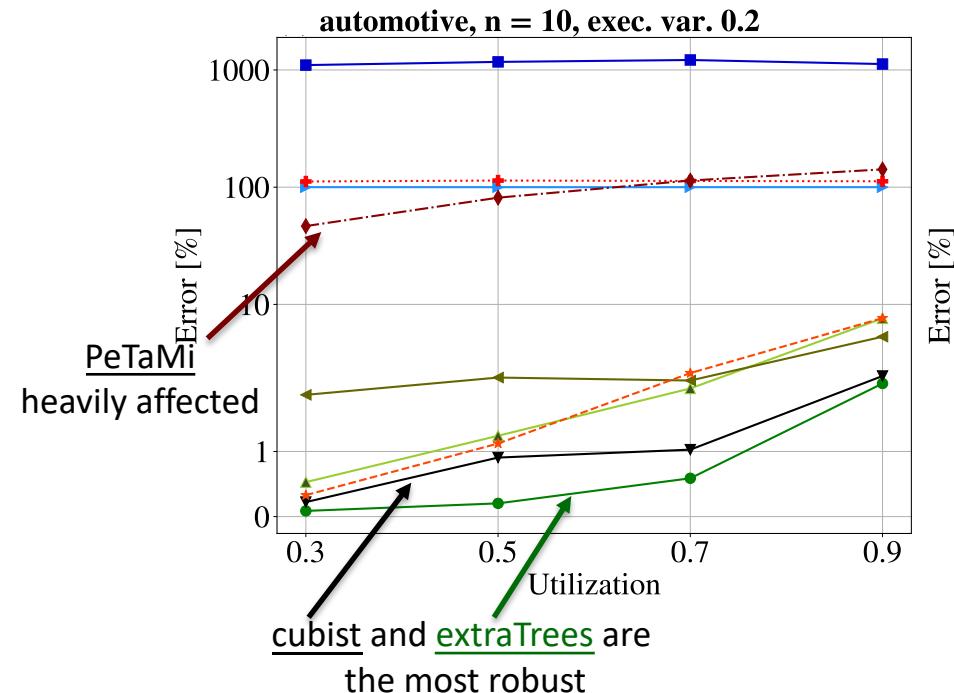
- extraTrees (green circle)
- cubist (black triangle)
- gbm (green triangle)
- bartMachine (brown triangle)
- avNNet (blue arrow)
- svr (blue square)
- Periodogram (red dashed star)
- Autocorrelation (red dotted star)
- PeTaMi (red diamond)



Robustness to variable execution time

extraTrees cubist gbm bartMachine avNNet svr
Periodogram Autocorrelation PeTaMi

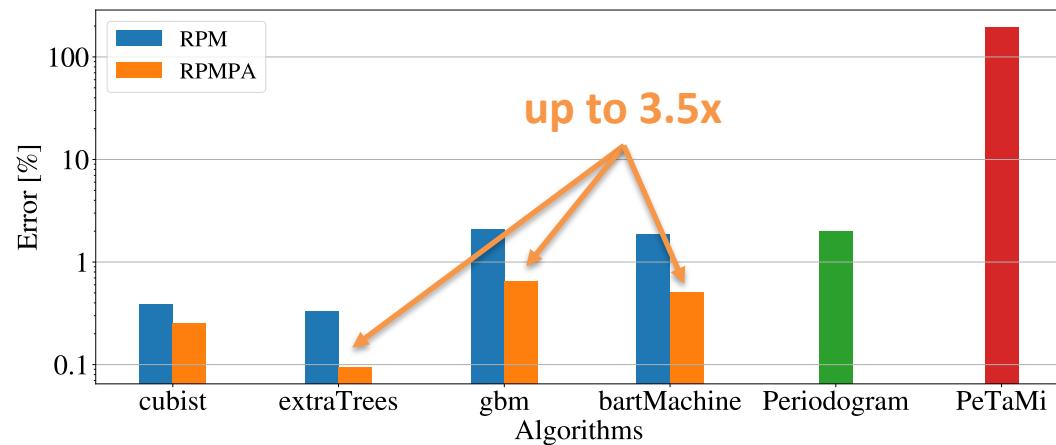
execution $\sim [(1 - \text{exec. var.}) \times \text{WCET}, \text{WCET}]$



Robustness to high-priority aperiodic tasks

- 12 automotive tasks
 - 6 periodic
 - 6 sporadic
- aperiodic tasks following a Poisson process suspending the execution
- task under analysis with medium priority

RPMPA reduces the error for all algorithms

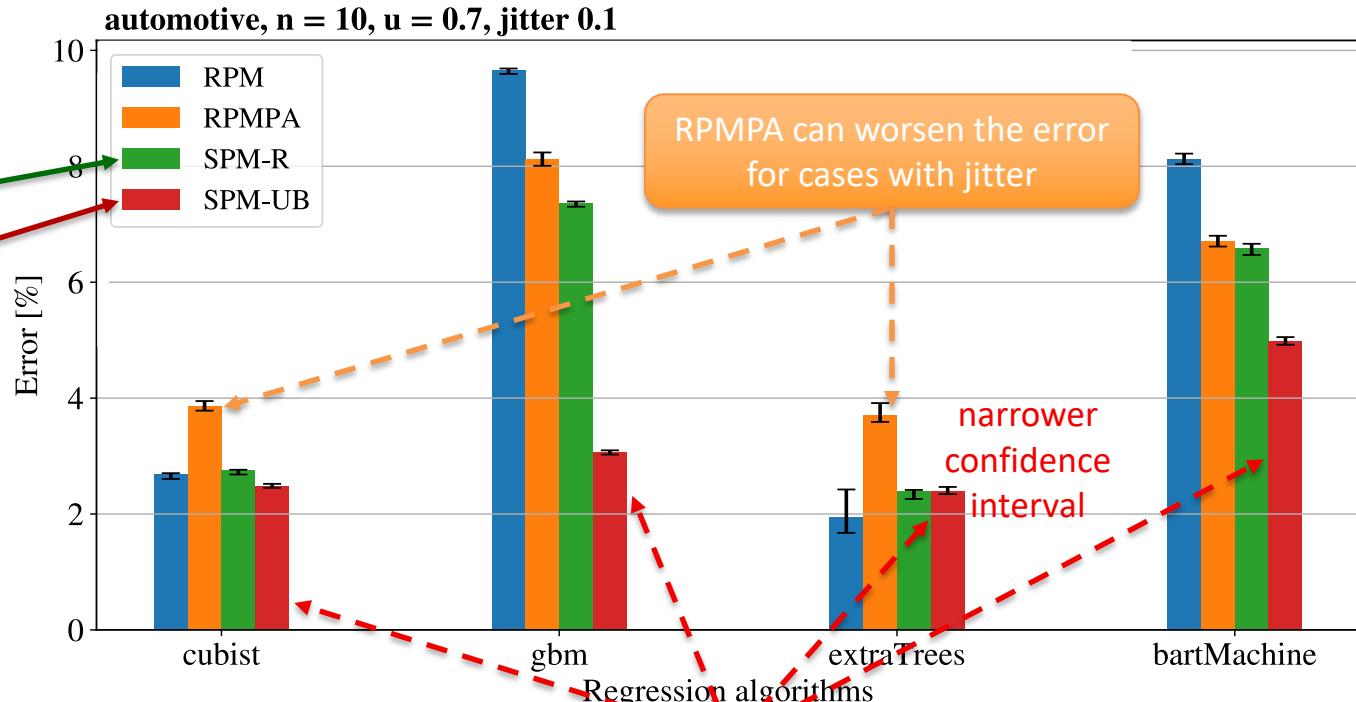


Space pruning on jitter

Q: What if no candidate is within the bounds?

A: We choose the prediction from:

- RPM
- upper bound



RPM – regression-based period miner

RPMPA – regression-based period miner with period adjustment

SPM – space pruning method

What is the performance on real data?

Two datasets [4] containing controller area network (CAN) messages obtained from vehicles

We have a lower error than the state of the art

RPM

RPMPA

Data set	Algorithm	RPM [%]	RPMPA [%]
CAN int.	cubist	3.2461	0.9703
	extraTrees	28.2568	1.6179
	com	15.8224	1.3919
	Machine	20.8588	14.0103
	TaMi	222.9416	-
	Periodogram	5.3368	-
CAN int.	cubist	5.5005	2.8963
		13.0256	1.7039
		19.8341	9.8881
		14.956	5.0264
	PeTaMi	177.6681	-
	Periodogram	9.5448	-

cubist has the **lowest** error

RPMPA significantly reduces the errors

How about efficiency?

Method	Average runtime [s]	
Candidate generation	5.101	
RPM (cubist)	2×10^{-6}	
PeTaMi	10.731	

RPM 2x faster than PeTaMi

RPMPA 2x faster than PeTaMi

SPM comparable to PeTaMi

Candidate selection – $O(m)$
 m – size of candidate list (40)

Candidate generation – $O(n \log n)$
Deriving bounds – $O(n)$

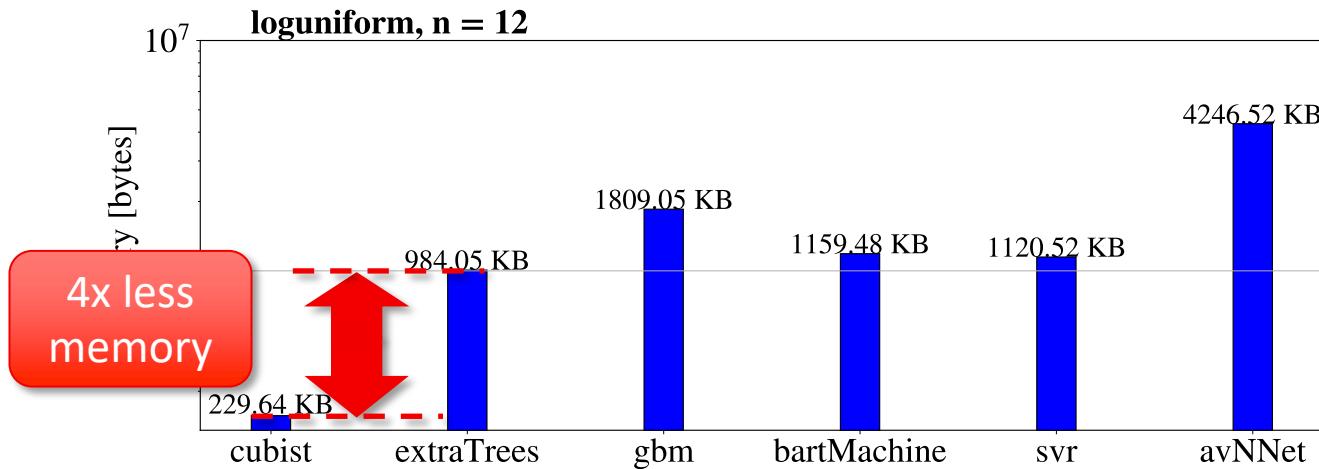
```
graph TD; subgraph Legend; end; subgraph Legend; end; subgraph Legend; end;
```

RPM – regression-based period miner

RPMPA – regression-based period miner with period adjustment

SPM – space pruning method

Memory comparison



Conclusions

1. *Can we use regression-based solutions for the period inference problem?*

- the first RBML method for period inference
 - two to three orders of magnitude **more accurate** than PeTaMi
 - **half the runtime** of PeTaMi

RBML – regression-based machine learning

O. Ilegorov, R. Torres, and S. Fischmeister. “Periodic task mining in embedded system traces”. *Real-Time and Embedded Technology and Applications Symposium*, pp. 331-340, 2017.

Conclusions

2. Which regression algorithm would result in a better accuracy for inferring periods?

- six regression algorithms (cubist, gbm, extraTrees, bartMachine, svr, and avNNet):
 - *cubist* has:
 - **the lowest** memory requirements;
 - **the lowest** runtime;
 - **the lowest** error on real traces.

gbm – gradient boosting

extraTrees – extremely randomized regression trees

bartMachine – Bayesian additive regression trees

svr – support vector regression

avNNet – averaged neural networks

Conclusions

3. Can we derive a pair of bounds to restrict candidate values?

- **lower bound** and **upper bound** on period values (space-pruning method).
 - lowered the error of **RPM** on traces with jitter up to **45%**.

RPM – regression-based period miner

Conclusions

4. How robust are our solutions to sources of non-determinism?

- robustness w.r.t. 
 - release time jitter ✓
 - execution time variation ✓
 - presence of high-priority aperiodic tasks ✓
 - missing jobs ✓
 - tardiness ✗
 - scheduling policy ✓
- *cubist* and *extraTrees* were **robust** against most types of interference.

Future work

1. Include information about higher-priority tasks running intervals to derive a **tighter** upper bound.
2. Extending the framework to support **multiprocessor** platforms.
3. Derive the periods of a task **without** requiring an **execution trace**.

Thank you!

